



❖ The Moving Average Filter — Explained from First Principles

🎯 Objective:

To reduce **noise** in a signal using a **simple low-pass filter** — specifically, a **moving average filter**.

A **moving average filter** works by replacing each sample in a signal with the **average of its neighbors**.

Think of it like:

"Instead of trusting one value, I'll look at a few values around it and smooth things out."
This is great for:

- Smoothing out short-term noise
- Keeping the overall shape of the signal intact
- Acting as a **low-pass filter** (because it removes rapid changes = high frequency)

🔍 What's Under the Hood?

Let's say we have this tiny signal:

Signal: [2, 4, 6, 8, 10]

If we apply a 3-point moving average filter, we use this kernel:

Kernel: [1/3, 1/3, 1/3]

Then each new value is:

New[1] = (2 + 4 + 6)/3 = 4

New[2] = (4 + 6 + 8)/3 = 6

New[3] = (6 + 8 + 10)/3 = 8

So the output is smoothed:

Filtered: [4, 6, 8]

🔗 This Operation Is Called **Convolution**

We "slide" the kernel over the signal and compute dot products at each step. This is the core of **convolution filtering**.

In Python, we use:

```
np.convolve(signal, kernel, mode='same'):
```

* `signal` : your noisy signal

* `kernel` : the filter (moving average here)

* `mode='same'` : keeps the output size equal to the input size

Why 'valid' mode?

*'valid': Keeps only the parts where the kernel fully overlaps (result is shorter).

*'same': Pads to maintain original length (useful for real-time filtering).

*'full': Even includes the parts with partial overlap (longer result).